

Club Cheeky — Prohibited Activities Protocols

- > Submission document for platform review. Each numbered section stands alone
- > and may be pasted into the corresponding question. Every control described
- > here is enforced in code (RLS, server-side RPCs, triggers, verification
- > state) — not aspirational policy. Cross-references:
- > `docs/Governance/aup-enforcement.md`, `community-safety.md`,
- > `verification-age.md`, `privacy.md`, `data-retention-deletion.md`.

The foundation — applies to all four below

Club Cheeky is a dating and matchmaking platform built on a verified-identity foundation. Protocols are layered — prevent, detect, respond, monitor.

- **Verification-as-entry (the door):** every member passes a government ID + live selfie check via Stripe Identity before they can interact with other members. Guests are confined to the public street level — they cannot message, match, or attend events. The door enforces an 18+ gate.
- **Identity anchoring:** one account per verified identity; gender is declared at signup; a mutual-compatibility filter means members only see and are seen by people within their stated preference.
- **Automated tripwires:** honeypot traps flag and shut down bot accounts at the database level (no messages, likes, waves, or event entries). All interactions run through server-side functions enforcing rate limits, blocks, and account status under row-level security. The client is never trusted. Messaging has hard daily caps with no paid bypass.
- **An AI staff on every floor:** the club is staffed by AI characters — a bouncer, a DJ, a bartender, a waitress, a hostess, and a general manager — present on every floor and privy to real-time member context. They engage members proactively and are trained to recognize signals of coercion, distress, or prohibited conduct. When a member raises any safety concern, the AI immediately directs them to the in-app Report/Block tools and reassures them a human reviews every report. Anything beyond the AI's authority escalates to the human owner for review, from any floor.
- **Response is human:** reports are reviewed by humans; we do not auto-ban on word filters. Violations escalate warning → temporary timeout → permanent bounce, with serious or illegal content skipping to a permanent bounce. Escalations are private, every ban is appealable via a human desk, and a dedicated safety email (`date.safely@smartscott.online`) is a direct line for concerns.
- **Monitoring & minimization:** Sentry and PostHog provide runtime and product observability; an owner back door gives the operator direct review of flags and account states. Raw ID material is never stored (Stripe Identity returns only a result), PII is segregated behind row-level security, messages are retained 3–90 days at the member's choice, and accounts can be deleted in-app.

1. Protocols for identifying and preventing Human or Sex Trafficking

****Prevention.**** The platform is structurally hostile to trafficking because it cannot be used anonymously. Every member is a verified adult: a government ID and live selfie checked through Stripe Identity, with an enforced 18+ gate and one account per verified identity. A trafficker cannot easily operate behind a real, checked identity, and a victim cannot be trafficked through a platform where every account is tied to a government ID.

****Detection.**** Beyond the door: honeypot traps remove bot/fraud accounts at the database level; server-side rate limits and row-level security bound abuse. The AI staff are present on every floor, privy to real-time member context, and trained to recognize coercion indicators — a member expressing they are being forced, controlled, surveilled, or unable to leave freely. The AI responds by directing the member to the in-app Report/Block tools and the safety desk, and escalates to the human owner.

****Response.**** Reports involving possible trafficking or coercion are prioritized and reviewed by a human. Where indicators warrant, we cooperate with law enforcement per our law-enforcement protocol — including account preservation on valid request and disclosure only under valid legal process. The safety desk (`date.safely@smartscott.online`) provides a confidential, direct line.

****Honest note.**** We do not scan message content automatically. Detection is behavioral (verification, honeypots, limits) plus human review of reports. Our posture is that the platform makes trafficking structurally difficult — not that it can detect it without member reporting or lawful process.

2. Protocols for identifying and preventing Illegal Acts

****Prevention.**** Our Terms and Acceptable Use Policy prohibit illegal conduct — harassment, threats, stalking, explicit content, hate, impersonation, illegal solicitation, and spam — with enforcement backed in code. The only payment rail is Stripe (memberships and in-app tokens); there is no marketplace for illicit goods, no peer-to-peer payments, and tokens are a server-side ledger only, never redeemable for cash.

****Detection.**** Bot and fraud accounts are caught by honeypot traps and shut down at the database level. Every interaction flows through server-side functions enforcing rate limits, blocks, and account status under row-level security. Sentry and PostHog provide observability, and the owner back door gives the operator direct review of flags and account states. The AI staff on every floor are privy to real-time member context and steer any disclosure of illegal activity into the human reporting process.

****Response.**** Reports are reviewed by humans. Illegal content results in a permanent bounce (no warnings for threats or illegal material), private escalation, and appeal through a human desk. Where lawful process demands, we cooperate with law enforcement per our law-enforcement protocol and preserve accounts on valid request.

****Honest note.**** We do not run automated content scanning of messages or photos. Detection is behavioral plus human review, and enforcement is escalation + lawful cooperation — we do not overclaim automated content understanding.

3. Protocols for identifying and preventing Impersonation

****Prevention.**** Impersonation is structurally prevented by verification: every account is bound to a government ID + live selfie checked through Stripe Identity (which includes face-matching the selfie to the ID). One account per person, one identity per account — fake profiles are a bouncing offense under our Terms, and anonymous sock-puppeting is not possible past the door.

****The platform's own AI never impersonates.**** The crew (Brutus, the DJ, Roxy, Trixie, Valentina, Chaz) are AI characters, clearly labeled as AI on every surface where they appear. Their house rules explicitly forbid claiming to be human or a real member. The platform itself does not engage in impersonation.

****Detection & response.**** Any report of impersonation — using someone else's photos, name, or identity — is reviewed by a human and results in a permanent bounce when confirmed. Blocked accounts cannot pair, message, or join events. Appeals are handled by a human desk.

****Honest note.**** Profile photos are not continuously re-verified after the door; impersonation via photos is caught by member reporting and human review, then removed permanently.

4. Protocols for identifying and preventing Child Sexual Abuse Material (CSAM)

****Prevention — the strongest layer.**** Club Cheeky is structurally adults-only: every member must present a government ID and live selfie through Stripe Identity, which enforces an 18+ gate that minors cannot pass. There is no anonymous or unverified route to the platform's interactive features. Explicit content of any kind is prohibited by our Terms and Acceptable Use Policy.

****Detection.**** Members can report from any chat in one tap. Reports are reviewed by humans — we do not auto-adjudicate on word filters. The AI staff, present on every floor and privy to member context, are instructed that any disclosure involving a minor or exploitative content goes immediately to the human process.

****Response — zero tolerance.**** Any CSAM report or discovery results in: immediate permanent account bounce; retention of relevant records in accordance with legal obligations; and mandatory reporting to the National

Center for Missing & Exploited Children (NCMEC CyberTipline) and law enforcement, in line with our law-enforcement protocol. We cooperate fully with valid legal process and preserve accounts on request.

****Honest note.**** We do not yet run automated perceptual hash scanning (for example, PhotoDNA-style matching) of uploaded images, and we will not overclaim that we do. Our CSAM posture is: minors cannot join (enforced ID verification), explicit content is prohibited and human-reviewed, and any report is a zero-tolerance permanent bounce with mandatory reporting to NCMEC and law enforcement.

5. Controls around profile verification during registration

Club Cheeky uses Stripe Identity itself to verify every member during registration — the platform's verification control is Stripe's own verification product.

- ****Verification is a prerequisite, not a checkbox.**** Registration creates a Guest confined to the public street level — no messaging, no matching, no events. Before any interactive feature, the member must pass the Door Check.
- ****The Door Check requires a government-issued ID and a live selfie,**** submitted through Stripe Identity, which validates the document, checks liveness, and face-matches the selfie to the ID. The check enforces our 18+ gate at the door.
- ****We store only the result**** — `verified\_at`, provider reference, and outcome. We never store the ID document, the selfie, or the ID number; raw verification material is not retained by us.
- ****Consent is explicit and separate:**** members consent to identity verification and biometric processing before the check runs, recorded in our `consents` table — never buried in the Terms.
- ****Identity is anchored:**** one account per verified identity; gender is declared at signup; the verification result powers the VIP badge, the free Silver card, and the welcome token grant.
- ****In-product framing:**** the Door Check is Brutus the bouncer's job — the AI staff direct members to the door and never bypass it.
- ****Failed checks grant nothing.**** Bot-flagged accounts cannot verify through to interactive features, and the operator can review verification state at any time through the owner back door.
- ****Minimization:**** account deletion wipes the verification state per our retention and deletion policy; a de-identified fraud flag may remain.

****Honest note.**** We rely on Stripe Identity's document/liveness checks as the verification authority; we do not run our own document-fraud analysis beyond what Stripe Identity provides.

6. Controls to prevent non-consensual involvement

Club Cheeky treats consent as the floor's law. Non-consensual involvement — including non-consensual intimate imagery or profiles created without a person's consent — is a permanent-bounce offense, and the platform is designed to make it structurally difficult and immediately reportable.

- **Verified identity:** every account is bound to a government ID + live selfie checked through Stripe Identity, so an account created without a person's consent is attributable to a real, verified identity — and is permanently removed on confirmation.
- **Profiles are self-created by verified members:** a profile is built by its own verified member with their own photos and declared identity. A profile created without consent is reported and removed with prejudice.
- **Report/Block from any chat, one tap:** non-consensual intimate imagery is explicitly covered by our Terms/AUP prohibitions on explicit content and sharing others' private information. Reports are reviewed by a human; confirmed cases are permanently bounced with no warning.
- **The AI staff on every floor** are privy to member context and trained to direct any disclosure of coercion, duress, or non-consensual content to the human reporting process and the safety desk.
- **Safety desk** (`date.safely@smartscott.online`) is a confidential direct line; reports involving potential victims are prioritized and, where warranted, referred to law enforcement per our protocols.
- **No-follow-up rule:** a decline closes the conversation — no re-messaging, no re-pairing in consecutive rounds, enforced in code.

Honest note. We do not run automated image-scanning for intimate imagery; detection relies on reporting, human review, and the verified-identity deterrent. Confirmed non-consensual content results in a permanent bounce and, where required, referral to law enforcement.

Living document — refined as the club grows. Backed by the policy layer in `docs/Governance/``. If automated content scanning is required by a platform or jurisdiction, it is a defined build we will ship and document here.*