

Takedown Notice & Appeals — Content Moderation Policy

- > **Status:** binding on the build (per the Governance README). Backs the
- > Stripe takedown-notice submission and the member-facing Acceptable Use
- > Policy. Where the code does not yet automate a step (content holds, banned
- > registry), the process below is the spec — operable today by human staff,
- > with automation to follow. Overclaims are the enemy: we describe what we
- > do and what we will build, never fiction.

1. How people report a policy-violating profile or content

Reporting is available through multiple, always-visible channels:

- **In-app Report/Block button** — one tap from any chat. The report is logged to our reports queue with the reporter, the subject, the content, and the timestamp.
- **The AI staff** — the crew (Brutus, the DJ, Roxy, Trixie, Valentina, Chaz) are present on every floor and take reports in conversation: they log the member, the incident, and the time, and provide assistance and routing if escalation is needed. They handle intake and triage; they do not adjudicate.
- **Direct support and safety contacts** — public mailboxes fielded by humans:
 - `info@smartscott.online` — general
 - `helpdesk@smartscott.online` — help and appeals
 - `date.safely@smartscott.online` — safety and reporting
 - `report-anonymous@smartscott.online` — anonymous reporting *(created; surfaced on the Contact page)*

2. How a submitted report is reviewed

- **Who:** review is layered. **Automated first pass:** a report is piped immediately to our automated image-review model — an OpenRouter-hosted reviewer that sits idle until a report triggers it (near-zero cost, fail-safe). It returns a verdict on the reported image/content within minutes (target: under 5). Clear violations are held and routed for human confirmation; inconclusive results go straight to human review. **Every ban is confirmed by a human.** The AI staff perform intake and triage in conversation; adjudication is human-led.
- **Immediate hold:** reported images/content are **held** (removed from public view) as soon as a report lands, while a decision is made — the safety of everyone involved comes before anyone's convenience. A hold is an action, not a verdict.
- **Targets (SLA):** automated first pass within minutes (under 5, aiming for 3); initial human review begins within a few hours of the report; straightforward cases resolve within 24–48 hours; escalated cases resolve within the appeal window below. Speed matters, accuracy matters more — a wrongful hold is lifted the moment the review clears it.
- **Actions — severity ladder:**
 - **5-year ban:** harassment, non-consensual behavior, and serious but non-criminal conduct. A report of non-consensual involvement lands a 5-year expulsion.

- **Permanent expulsion:** sexually explicit behavior — unsolicited explicit imagery or exposure. Gone for good, no exceptions.
- **Permanent + authorities:** anything that appears to violate state or federal law — reported to the local authorities, with evidence preserved for lawful request.

3. How a reported user appeals a removal or ban

- **Window:** the affected member has **14 days** from the notice to appeal. If no appeal is filed within the 14-day limit, the decision is assumed correct and the punishment stands.
- **Human review:** a timely appeal triggers an exhaustive review by human staff. Both sides may submit evidence to substantiate their claims within a 14-day evidence period.
- **Decision:** at the end of the evidence period, a decision is made on the supporting evidence, with an interview arranged by the support desk (by phone or video call) if needed — within 3 days. The full appeal completes within roughly one month: expedient and thorough.
- **Wrongful bans are made right:** if a verdict of not guilty overturns a hold or ban, the hold is lifted, the account is restored, and any time lost on the platform is compensated per the member's choice — membership credit of equal value, or a refund through our payment processor.
- **Notice:** every outcome is communicated by the support desk by email — the decision, the examination results, and, where applicable, a detailed reason for how the conclusion was reached.

4. Banned-account registry & retention

- **Registry:** expelled accounts are flagged and recorded in a banned-account registry, consulted at membership/signup so banned users cannot simply re-register. Repeat offenders returning after serving a ban are banned for the life of the user.
- **Retention:** for incidents that may violate state or federal law, the username and non-government information are retained, along with any evidence collected during the investigation, for the length of the ban (5 years) in the event authorities request it. Standard records follow the data-retention policy (`data-retention-deletion.md`).

5. Cross-references

Doc	Relationship
``aup-enforcement.md``	Detection, escalation ladder, report/block mechanics
``community-safety.md``	Conduct rules and the human-review process
``terms.md``	Bouncing, conduct, liability
``data-retention-deletion.md``	Retention windows and deletion
``stripe-prohibited-activities.md``	The platform-level prohibited-activity submission

Submission answer (final, copy-paste)

****Reporting.**** Members can report a policy-violating profile or content through three channels, all fielded by humans: the in-app Report/Block button in any chat — one tap, logged with the member, the content, and the time; the AI staff, who take reports in conversation and route them to the review process; and the public support and safety contacts (info@smartscott.online, helpdesk@smartscott.online, date.safely@smartscott.online, and report-anonymous@smartscott.online).

****Review.**** A reported image or piece of content is taken down immediately and held until a decision is reached — a block is placed on the user's profile content in the database the moment the report lands. The content is reviewed by the AI staff, with a response in under five minutes. If the allegation is unfounded, the block is lifted and an apology email explains the incident and the procedure. If it is founded, the responsible user is immediately expelled from the app, and an email explains the incident, the findings, and how the conclusion was reached. The account is then flagged and moved to a banned-account registry that is checked whenever anyone applies for membership, so a banned user cannot simply re-register. If the incident appears to violate state or federal law, the user's username and non-government information are retained — along with any evidence collected during the investigation — for five years, the length of the ban, in the event authorities request them. Repeat offenders who return after serving a ban are banned for life.

****Appeals.**** Any user subject to a ban or punishment has 14 days to appeal. If no appeal is filed within 14 days, the decision stands. A timely appeal triggers an exhaustive human review, and both sides may submit evidence during a 14-day period. At the end of that period, a decision is made on the evidence — with an interview conducted by video chat when needed — within three days, completing the appeal within one month: expedient and thorough. If the verdict is not guilty, the account is restored and any time lost is compensated at the member's choice: a refund to the original credit or debit payment method, or membership of equal value.